# Moving Cost Magnitudes in Moving Cost Models

Greg Howard*

December 20, 2023

**Abstract**

What is the correct interpretation of average moving costs? I show that in the steady-state of a standard moving cost model, average moving costs are proportional to the difference between the Shannon entropy of next period's location and the Shannon information of staying in the same location. Therefore, moving costs are correctly interpreted as a statistic about the modeler's lack of information regarding future moves, but not about the actual cost of moving in the real world. This alternative interpretation helps make sense of the wide range of moving costs in the literature.

Many papers in economics estimate the average cost of moving to be quite large, often several times annual household income (Kennan and Walker, 2011; Bryan and Morten, 2019).[1] This may seem implausibly large when compared to actual expenses associated with a move. However, others have noted that these migration costs could reflect other frictions, too, and that if they explicitly model these other frictions, then estimated moving costs fall (Schmutz and Sidibé, 2019; Porcher, 2020; Heise and Porzio, 2022; Giannone, Li, Paixao and Pang, 2023).[2]

Jia, Molloy, Smith and Wozniak (2023), a review article in the *Journal of Economic Literature*, summarizes the state of the literature as, "while unobserved and potentially very large costs might help explain migration rates that are low relative to the potential earnings gains from migration, different models imply substantively different estimates of the size of these costs."[3]

In this paper, I propose a different way to think about these estimated moving costs. I show that average moving costs are a measure of the modeler's information about future locations. A corollary to this interpretation is that estimated moving costs depend on arbitrary decisions of the modeler, such as the time period or the geographic partition. This corollary rules out a literal

---

[1]In Table 2, I show a range of large moving cost estimates in the literature.

[2]Perhaps because of this debate, the literature that uses moving costs in their models is much larger than the number of papers that report moving costs as a main outcome. For example, Caliendo, Dvorkin and Parro (2019) has a similar model that would imply similarly sized moving costs, but develops solution techniques that do not require backing out the moving cost parameters. Schubert (2021) does not report the average moving cost, but does consider counterfactuals in which the moving costs change.

[3]Other methodologies of uncovering migration costs also give various different results. Koşar, Ransom and van der Klaauw (2022) uses a survey to the estimate the willingness to pay to avoid moving, and estimates an average moving cost of $54,000. Bryan, Chowdhury and Mobarak (2014) induces 22% of rural Indonesian workers into seasonal migration with an incentive of $8.50.

interpretation of moving costs and makes the debate about the size of these costs somewhat moot.

To establish my interpretation, I show algebraically that in the steady-state of a standard moving cost model, moving costs are proportional to the average Shannon entropy of next period's location minus the Shannon information of next period's location being the same as the current location.[4] In other words, moving costs are a measure of information: specifically, how surprised the modeler is by the agent's future location compared to the agent not moving.

Based on this result, I show that estimated moving costs change when the time period considered by the model is different, or when the model partitions geographies in a different manner. A more novel result is that they are also sensitive to the modeler's information about the agents. For example, knowing the birthplace of each person leads the modeler to estimate smaller moving costs. I give examples of the ways these arbitrary decisions affect moving costs using data from the 2000 Census and the American Community Survey. Because none of these modeling choices affects the actual decisions of people on the margin of moving, I argue that interpreting moving costs literally is a mistake.

However, that does not mean moving costs are uninteresting. Given the formulae I provide, one can interpret moving costs as a measure of the information of the modeler. So comparing moving costs across models is informative of how good those models are at predicting future locations. This alternative interpretation makes sense of some recent results, specifically that richer mod-

---

[4]The concepts of Shannon entropy and Shannon information were developed in Shannon (1948).

els of moving—which typically incorporate more information—exhibit smaller moving costs (Zerecero, 2021; Giannone et al., 2023; Heise and Porzio, 2022; Porcher, 2020; Schmutz and Sidibé, 2019).

Of course, if the standard moving cost model were the true model of the world, then moving costs could be interpreted both in the way I describe here and also as a literal average moving cost. Therefore, my argument that migration costs should not be taken literally implies that the standard moving cost model is misspecified. To conclude the paper, I focus on the i.i.d. preference shocks in the model and argue that these are the source of the misspecification. Relaxing this assumption may allow future work to better estimate average moving costs.

Besides the literature that uses moving costs in their models, which I will discuss in detail in Section 3, this paper also has some similarities to the literature that relates discrete choice models to generalized entropy (e.g. Jose, Nau and Winkler, 2008; Fosgerau and de Palma, 2016). These papers argue there is an equivalence between utility maximization and entropy minimization in discrete choice models. To my knowledge, while this literature describes many interesting relationships between utility maximization and entropy, no one has related the estimated moving costs to entropy as I do here.[5] The key assumption that allows me to show my interpretation of moving costs is that

---

[5]Porcher (2020) is perhaps the closest paper to this one, in that it has to do with both Shannon entropy and internal migration. That paper assumes rationally inattentive agents, and a typical result in rational inattention is that the costs that agents have to pay is related to the Shannon entropy of the information they acquire. This is equivalent to a discrete choice problem (Matějka and McKay, 2015). However, there is a huge difference between Porcher (2020) and this paper because this paper emphasizes the moving costs as a measure of the modeler's lack of information, whereas that paper emphasizes that agents' lack of information can look like moving costs.

4

I focus on a setting with a steady-state, where changes in the utility net of preference shocks and moving costs will cancel out.

# 1 Standard moving cost model

In this section, I use the standard moving cost model to derive an interpretable expression for average steady-state moving costs. In this model, agents choose their location to maximize the present value of their utility. As part of that utility, they face moving costs and they draw i.i.d. extreme value shocks for each location each period that induce some of them to move and some of them to stay.[6]

There is a continuum of people indexed by $n$ that live in discrete locations indexed by $i$. Time is also discrete and is indexed by $t$. The population of people living in $i$ at time $t$ is denoted by $p_{it}$. The share of people in $i$ who move from $i$ to $j$ at time $t$ is denoted $m_{i \to j,t}$.[7] $m_{it}$ denotes the total outmigration share from $i$ to all locations $j \neq i$ at time $t$. When referring to steady-states, the $t$ index is dropped. Moving costs are bilateral between two locations, so $\delta_{ij}$ refers to the moving cost from $i$ to $j$. I assume there is no cost to not moving, i.e. $\delta_{ii} = 0$. I will use the notation $\mathbb{E}_i$ to refer to the population-weighted average across locations. I will use the notation $\mathbb{E}^m$ to refer to the migration-

---

[6]Some versions of the standard model include the i.i.d. extreme value shocks as part of the moving costs. For example, Kennan and Walker (2011) models moving costs this way. They report average moving costs for all possible moves, but report moving costs paid by actual movers that are much lower, in fact negative. However, their most well-known statistic is measured similarly to the statistic I consider below: "For the average mover, the cost is about \$312,000 (in 2010 dollars) if the payoff shocks are ignored" (Kennan and Walker, 2011, p. 232).

[7]Based on this notation, $m_{i \to i,t}$ will refer to the non-migration rate in $i$.

weighted average. I will be particularly interested in the average migration cost, which we define to be $\bar{\delta} \equiv \mathbb{E}^m[\delta_{ij}]$.

In this section, I assume agents are homogeneous except for their location. I also only consider average moving costs in the steady-state of the model. I relax both of these assumption in Appendix A.

Agents maximize the present value of utility, represented by this value function:

$$V_{nt}(i) = \max_j \log w_{jt} + a_{jt} - \delta_{ij} + \frac{1}{\mu}\epsilon_{jnt} + \beta\mathbb{E}V_{nt+1}(j)$$

where $w_{jt}$ is the (real) wage, $a_{jt}$ is the amenities in $j$, $\delta_{ij}$ is the moving cost from $i$ to $j$, $\epsilon_{jnt}$ is a time-person-location i.i.d. extreme value shock. $\beta$ is the discount factor, and $\mu$ is a scale parameter, which governs the elasticity of substitution between places.

Define $v_{jt} \equiv \log w_{jt} + a_{jt} + \beta\mathbb{E}V_{nt+1}(j)$. Then the migration rate is given by

$$m_{i\to j,t} = \frac{\exp(\mu(v_{jt} - \delta_{ij}))}{\sum_k \exp(\mu(v_{kt} - \delta_{ik}))}$$

Because $\delta_{ii}$ is normalized to zero, $\delta_{ij}$ can be shown to equal:

$$\delta_{ij} = v_{jt} - v_{it} - \frac{1}{\mu}\log m_{i\to j,t} + \frac{1}{\mu}\log m_{i\to i,t} \tag{1}$$

Intuitively, moving costs are decreasing in migration. In this paper, I wish to focus on the following statistic which is often reported in papers in the literature, the migration-weighted average moving cost in the steady-state of

the model:

$$\bar{\delta} \equiv \mathbb{E}^m[\delta_{ij}] = \frac{\sum_{i,j:i \neq j} p_i m_{i \to j} \delta_{ij}}{\sum_{i,j:i \neq j} p_i m_{i \to j}}$$

The main proposition relates $\bar{\delta}$ to measures of information about future locations. Before I write down the proposition, it is helpful to define some notation.

Define $J$ to be a discrete random variable, which is the next period's location. Lower-case $j$ will refer to specific realizations of $J$. I will use the notation $H(J|i)$ to refer to the Shannon entropy of $J$ for a person currently living in $i$, and the notation $I(j|i)$ to refer to the Shannon information of the realization of $J = j$ (Shannon, 1948). Since $m_{i \to j}$ is the migration probability for someone living in $i$ to move to $j$, this means that

$$I(j|i) = -\log m_{i \to j}$$

and

$$H(J|i) = -\sum_j m_{i \to j} \log m_{i \to j}$$

**Proposition 1.** *In the steady-state of the standard moving cost model, the average moving cost is the average Shannon entropy of next period's location minus the Shannon information of next period's location being the same as this period's location, all divided by the average moving rate times the migration*

7

*elasticity. In math,*

$$\bar{\delta} = \frac{1}{\mathbb{E}_i m_i} \frac{1}{\mu} \mathbb{E}_i[H(J|i) - I(i|i)] \qquad (2)$$

*Proof:* Plugging in (1) to the definition of $\bar{\delta}$,

$$\bar{\delta} = \frac{1}{1 - \sum_i p_i m_{i \to i}} \sum_{i,j:i \neq j} p_i m_{i \to j} \left( -\frac{1}{\mu} \log m_{i \to j} + \frac{1}{\mu} \log m_{i \to i} \right)$$

Note that in steady-state, the $v_{it}$ and the $v_{jt}$'s all cancel out because $\sum_k p_k m_{k \to i} = \sum_k p_i m_{i \to k}$. The steady-state assumption is a key assumption that allows my interpretation, and which is why my results do not extend to more-general discrete choice models. Rearranging,[8]

$$\bar{\delta} = \frac{1}{1 - \sum_i p_i m_{i \to i}} \frac{1}{\mu} \sum_i p_i \left( -\sum_j [m_{i \to j} \log m_{i \to j}] + \log m_{i \to i} \right)$$

Recall that I defined $m_i \equiv \sum_{j:j \neq i} m_{i \to j}$ to be the total outmigration from $i$. Using the definitions of Shannon entropy and Shannon information,

$$\bar{\delta} = \frac{1}{\mathbb{E}_i m_i} \frac{1}{\mu} \mathbb{E}_i[H(J|i) - I(i|i)]$$

$\square$

An informal way to understand Shannon information is that it measures how surprising an event is. Since most people do not move, the event of not

---

[8]To derive this expression from the one above, note that $\sum_{j \neq i} m_{i \to j} \log m_{i \to 1} = (1 - m_{i \to i}) \log m_{i \to i}$. The $-m_{i \to i} \log m_{i \to i}$ term is then moved into the other term so that the sum is over all $j$, and not just $j \neq i$. This leaves the $\log m_{i \to i}$ term outside the summation.

moving is unsurprising, and the Shannon information of not moving is small. Shannon entropy measures the expected Shannon information. So if it is very hard to predict where people will live next period, then the Shannon entropy will be large.

Another way to think about Shannon entropy is that Shannon entropy is approximately proportional to the number of "yes or no" questions one would have to ask in order to acquire the information, i.e. the number of bits the information contains.[9] So $H(j|i)$ is proportional to the bits of information that you would need to communicate where a person currently in $i$ will choose to live next.

In addition to the previous result of equation (2), I can alternatively relate the average moving cost to the Shannon entropy of future locations *conditional on migrating*.

Define an event $m$ to be when the random variable $J$ takes on any realization that is not $i$, i.e. a move. I will use the notation $H(J|i, m)$ to be the conditional Shannon entropy of next period's location given that the agent moves.

**Proposition 2.** *In the steady-state of the standard moving cost model, average moving costs are the migration-weighted average Shannon entropy of next period's location conditional on moving plus the Shannon information of the moving minus the Shannon information of not moving, all divided by the mi-*

---

[9]This is an approximation because Shannon entropy is a continuous measure. It also has to be scaled by $\log 2$ to convert the units of Shannon entropy into bits.

*gration elasticity. In math,*

$$\bar{\delta} = \frac{1}{\mu}\mathbb{E}^m\left[H(J|i,m) + I(m|i) - I(i|i)\right] \qquad (3)$$

*Proof*: Define $m_{i\to j}^* = \frac{m_{i\to j}}{m_i}$ to be the probability of moving to $j$, conditional on moving at all. Then, we can algebraically rearrange the expression for average moving costs as:

$$\bar{\delta} = \frac{1}{\sum_i p_i m_i}\frac{1}{\mu}\sum_i p_i m_i \left(-\sum_{j\neq i}[m_{i\to j}^* \log m_{i\to j}^*] - \log m_i + \log(1 - m_i)\right)$$

So

$$\bar{\delta} = \frac{1}{\mu}\mathbb{E}^m\left[H(J|i,m) + I(m|i) - I(i|i)\right]$$

$\square$

The two key assumptions in this setup are that the model is in steady-state, a condition that is never exactly true in the data, and that the agents are homogenous but for their initial location. In Appendix A, I show that if either of these assumptions is relaxed, then the expressions for $\bar{\delta}$ are very similar to the equations (2) and (3), but with an additional term that represents the average gain from migration net of moving costs and idiosynchratic shocks.[10]

These propositions have three important implications for interpreting migration costs, which I cover in these three corollaries.

---

[10]I can quantify how important this term is. Relaxing just the steady-state assumption, and using the same data from Section 2, this additional term is quantitatively negligible, about 0.4 percent the size of the terms in equations (2) or (3). Intuitively, one of the reasons this term is so small is because gross migration is much larger than net migration (Jia et al., 2023), so the data is not too far away from steady-state.

**Corollary 1.** *Average moving costs depend on the modeler's choice of length of the time period.*

Over short time horizons, the Shannon entropy, conditional on moving, does not vary much. Whether I look at 1 year or 5 years, the percentage of migrants from state $i$ moving to state $j$, $m^*_{i \to j}$ is roughly constant. However, migration rates are smaller for shorter time horizons. So based on equation (3), average moving costs will vary with the time period chosen as the latter terms, $I(m|i) - I(i|i) = \log \frac{1-m_i}{m_i}$, will increase when time horizons are short. In fact, as time horizons get arbitrarily short, estimated average moving costs will get arbitrarily large.[11]

**Corollary 2.** *Average moving costs depend on the modeler's choice of geographic partition.*

The Shannon entropy of next period's location depends on how the modeler partitions geography. Generally, the more locations there are, the harder it is to predict exactly which one any given person will end up in. Therefore, one would expect that Shannon entropy would increase in the number of locations.[12] Mechanically, migration rates also increase in the number of locations. As far as I know, there is no way to order geographies such that estimated migration costs must increase or decrease, but in the empirical results, I show

---

[11]When we express migration costs in utility terms, one might be worried that the correct way to interpret those costs also depend on the time horizon. For example, when I convert to dollar terms, and I consider five-year time periods, we will calibrate $w$ to five times the yearly wage, but when I consider one-year time periods, I will calibrate $w$ to the yearly wage. However, unless migration over time horizon $t$ is proportional to $\exp(t)$, this will not cancel. In the data, it is closer to proportional to $\sqrt{t}$ (Howard and Shao, 2022).

[12]This does not have to be true in all cases, if the more precise location is extremely informative of which locations the person is likely to move to.

that the Shannon entropy change dominates the change in the information of not moving when I apply it to states versus migration public use microdata areas (MIGPUMAs). Certainly, there is no reason to expect the change in Shannon entropy and the change in migration rates to cancel out.

In fact, let us consider two silly examples to show that moving costs can be estimated to be very large or very small depending on the partition. In the first case, consider partitioning every house into its own geography. In the 2000 Census, 43 percent of people moved houses in the previous 5 years. But the Shannon entropy conditional on moving is enormous because it is almost impossible to predict the exact house that anyone would live in. So based on equation (3), we would have an enormous number plus log(57/43). Just to put a number on it, we can assume that modeler can assign no individual house a probability of being chosen of greater than 0.1 percent. Then a lower bound on this enormous number would be $-\log\frac{1}{1000} + \log\frac{57}{43} \approx 7.2$.

Alternatively, we could partition the United States into houses with an even-numbered address and ones with an odd-numbered address. If we assume that it is random which type of house you move into, we would expect 21.5 percent of the population to be "moving regions" each year. But conditional on moving, the Shannon entropy is zero. So the estimated average moving cost would be $\log\frac{78.5}{21.5} \approx 1.3$.

**Corollary 3.** *Average moving costs depend on the modeler's information set.*

Suppose the modeler knew some immutable characteristic about people, denoted by $s$, such as their race or their birthplace. If they estimate separate

moving costs by this characteristic, then equation (2) becomes[13]

$$\bar{\delta}_s = \frac{1}{\mathbb{E}_{is} m_{is}} \frac{1}{\mu} \mathbb{E}_{is} \left[ H(J|i, s) - I(i|i, s) \right] \tag{4}$$

Shannon entropy is convex, and Shannon information is concave, so by Jensen's inequality, $\mathbb{E}_{is}[H(J|i, s) - I(i|i, s)] \leq \mathbb{E}_i H(J|i) - I(i|i)]$. Since $\mathbb{E}_{is} m_{is} = \mathbb{E}_i m_i$, the whole quantity, $\bar{\delta}_s$ is weakly smaller that $\bar{\delta}$. And if $s$ provides any information about the next periods' location, then the inequality will be strict.

This is also true for some characteristics of a person that are not immutable. For example, if the modeler modeled the decision making process in two stages where, first, each person came up with a consideration set, and second, compared the utilities available in each, $s$ could be the consideration set.[14] In this setup, I can still derive formula (4).[15] So in a model with consideration sets, the modeler will estimate lower moving costs than in a model without consideration sets.

In the limit, if the modeler could perfectly predict everyone's next period location (e.g. $s = j$), then moving costs would be reduced to negative infinity. This is because not moving would be an infinitely large surprise for someone

---

[13]See Appendix A.3 for the derivation.

[14]As this example illustrates, the characteristic $s$ does not need to be measured in the data. It can be something the modeler can only see with the model.

[15]For the derivation, see Appendix A.4.

While immutable characteristics and consideration sets do not break the main results, this is not true of all possible characteristics. In Appendix A, I add a general state variable to each person, which can affect their migration costs and location utilities. In general, I show that migration costs are the sum of two terms: one which is the difference between the Shannon entropy of next period's location and the Shannon information of not moving, and one that represents the average gain from moving net of migration costs and idiosynchratic shocks. In these two examples, that second term is zero, but that does not need to generally be true.

they knew was moving.

# 2 Moving cost calibrations with data

In this section, I take the observations from the last section and illustrate them using real world data.

In particular, I will estimate the average moving costs using equation (2) with data, using the Census and the American Community Survey from 2000 (Ruggles, Flood, Sobek, Brockman, Cooper, Richards and Schouweiler, 2023).[16] For each state-pair, I calculate $m_{i\rightarrow j}$ as the share of people who lived in state $i$ that moved to state $j$, either from 1995 to 2000 in the Census, or from 1999 to 2000 in the ACS. I also calculate $m_{i\rightarrow j,b}$, where I calculate the probability of moving from $i$ to $j$ given a birthplace $b$. And I also calculate $m_{i\rightarrow j}$ where $i$ and $j$ are MIGPUMAs instead of states.[17]

Kennan and Walker (2011) and many subsequent papers often express moving costs in dollar terms instead of utility. Since wages are in logs in the model setup, one can interpret these average moving costs as a percent of wages.[18] Therefore, one might think of moving costs as a measure of the expected Shannon information minus the Shannon information in the event of not moving, where each bit of information "costs" $\frac{w}{\mu \log 2}$ dollars per migrant.

---

[16]This is the only year, to my knowledge in which similar surveys asked about the 1-year migration rate (the ACS) and the 5-year migration rate (the Census).

[17]MIGPUMA stands for Migration Public Use Microdata Area, and is a within-state region with at least 100,000 people.

[18]Kennan and Walker (2011) actually expresses utility in dollar terms directly, so there is no need for this adjustment. However, much of the subsequent literature does express wages in logs.

Table 1: Estimated Moving Costs for Different Models

|  | Shannon Entropy | Migration Rate | Estimated Moving Cost | Cost in $1000's |
|---|---|---|---|---|
| 1 year, states | 0.182 | 0.024 | 6.692 | 315 |
| 5 year, states | 0.561 | 0.085 | 5.585 | 1312 |
| 5 year, states (modeler knows birthplace) | 0.512 | 0.085 | 4.981 | 1171 |
| 5 year, MIGPUMAs | 1.231 | 0.173 | 5.983 | 1406 |

Notes: All datasets are from 2000. 1 year migration uses migration measured over 1 year from the ACS. 5 year migration uses migration measured over 5 years from the Census. The unit of geography is a state or a MIGPUMA, a subset of a state with at least 100,000 people in it. Birthplace is an indicator variable either for the state of birth or for being from anywhere outside the 51 U.S. states. The median household income in 2000 (for people also living in the U.S. in 1995) was $47,000, so for 1-year migration, the last column is that times the estimated moving cost. For 5-year migration, the last column is $47,000 times five times the estimated moving cost.

I then calibrate the average moving costs according to equation (2), assuming $\mu = 1$. In the literature, there is little consensus on what $\mu$ is, and some good arguments that typical methods have not estimated it well (Borusyak, Dix-Carneiro and Kovak, 2022), so I use $\mu = 1$ not because I believe that but because it is easy for the reader to scale the moving costs by whatever $\mu$ they prefer.[19] The comparisons of results are intuitive. In the 1 year calibration, I estimate moving costs of 6.7 log points, or when converted to dollars, $315,000. This is not too different from Kennan and Walker (2011), who estimated moving costs of $312,000 (p. 232).[20]

---

[19]Borusyak et al. (2022) makes the point that regressing the change in population on labor demand shocks—even well-identified labor demand shocks—does not identify $\mu$ because the shocks are typically highly correlated across space, and therefore affect both origin and destination locations for many migrants. They suggest a non-linear least squares estimation instead.

[20]The fact that these are only $3000 different is mostly a coincidence. Kennan and Walker (2011) is using 2010 dollars, while I use 2000 dollars, and the model in Kennan and Walker (2011) is much richer. They also explicitly model the elasticity of migration, which in their

In the 5 year calibration, I estimate smaller moving costs in utility terms because many more people are moving: 5.6 log points. In dollars, this number is larger—$1.31 million—because I multiply by five years of wages instead of one.

If the modeler knows the birthplace, the entropy decreases since birthplace is a helpful predictor of future location choices. Compared to the 5 year calibration where the modeler does not know birthplace, the moving cost is even lower: 5.0 log points ($1.17 million). This is consistent with Zerecero (2021) who finds lower estimated moving costs in a model with birthplace.

Finally, if I use MIGPUMAs instead of states, it is much harder to predict future locations, since MIGPUMAs are a finer geography. The moving costs increase by about 0.4 when I use MIGPUMAs instead of states, to 6.0 log points ($1.41 million).

# 3    Discussion

## 3.1    How should moving costs be interpreted?

Is there one of these moving costs that is more "correct" than the other ones? No. The differences depend on arbitrary choices made by the modeler. Even very natural choices, such as MIGPUMAs vs. states, 1-year vs. 5-year, and whether to include information on birthplace lead to large differences in estimated migration costs, of about half a log-point or more. As I showed in Section 1, if the modeler makes even less "normal" assumptions, the migration

---

paper is a semi-elasticity since they have linear utility in consumption.

costs could diverge to infinity or negative infinity.

Because these arbitrary decisions affect the estimated moving cost, it therefore cannot be a reasonable measure of the actual cost of moving.

So how should a reader interpret reported moving costs in an economics paper? They should not try to gauge if the moving cost is big or small compared to their prior on how expensive it is to move. It is also not that useful to compare it to the perceived benefits of moving, as many papers in the literature do.[21] However, they may want to compare the moving costs to other papers or other model specifications, as I do in Table 2.[22] These comparisons can give a sense of how much information the model has. The larger the average moving cost, the less the model is able to predict where people will be in the next period, relative to the information of staying in place.

---

[21]One caveat is if the model features a rich state space, with large average gains to moving net of migration costs and idiosynchratic shocks. When these terms are big compared to the information terms I focus on in this paper, this comparison is more apt.

[22]To include a paper in this table, I required the paper to report an average moving cost in some sort of interpretable units and to use extreme value shocks. Papers such as Bishop (2012) and Oswald (2019) report a moving cost function, and seem to have moving costs in the same ballpark as Kennan and Walker (2011), but do not report average costs. Bartik and Rinz (2018) reports an average moving cost of $683,000, but this is not the average of all movers; rather it is the average cost for a 500 mile move. Similarly, Bayer and Juessen (2012) also does not feature extreme value shocks, so the moving costs are not exactly a measure of information. Nonetheless, Bayer and Juessen (2012) does estimate substantially smaller moving costs ($34,248), likely because they incorporate information about migrants persistent preferences over locations.

| Paper | (1) Estimated Migration Costs | (2) Length of time | (3) Geography | (4) Modeler's Information | (5) Notes |
|---|---|---|---|---|---|
| Tombe and Zhu (2019) | 282% of lifetime income | lifetime | Chinese provinces × urban/rural | birthplace | Paper reports the parameter which I called $\bar{\delta}$ as 2.82 which I interpreted as a share of lifetime income because in their model, moving costs are paid every year a migrant is away from their birthplace |
| Zerecero (2021) | 56% of lifetime consumption | 1 year | French départements | current location and birthplace | Without home bias, migration costs estimated to be 10% larger |
| Bryan and Morten (2019) | 39% of lifetime income | lifetime | Indonesian regencies | birthplace | |
| Bryan and Morten (2019) | 15% of lifetime income | lifetime | U.S. States | birthplace | |
| Ransom (2022) | $394,000 to $459,000 (2004-2013 dollars) | 1 year | 35 U.S. core-based statistical areas | current location, work experience, age, employment and labor force status, and unobserved type | |
| Kennan and Walker (2011) | $312,146 (2010 dollars) | 1 year | U.S. States | current location, birthplace, current wage, age, type (stayer or mover), last year's location, and wage at that location | |
| Giannone et al. (2023) | 196,202 CAD (2016 dollars) | 2 years | Canadian provinces | current location, wealth, income shock, age, housing tenure status, and housing consumption | |
| Porcher (2020) | 75% of annual earnings | 1 year | Brazilian mesoregions | current location, information acquired by the agent about productivity in different locations | Without information frictions, migration costs estimated to be 40% larger |
| Heise and Porzio (2022) | 3.1%-5.3% of lifetime income | continuous | 4 German regions | current location, home location, current employment status, current wage, location of job offer, wage of job offer | While the model is continuous time, workers only consider moving at discrete times when they get a job offer |

Table 2: Migration Costs in the Literature

Table 2 is roughly ordered by the size of the moving costs, from largest to smallest.[23] In column (4), where the modeler's information is listed, the amount of things that the modeler knows increases as the moving cost decreases. Of course, the geographies and timeframes are changing as well, so that will also affect the moving costs, but the information column seems to be playing a large role.

For example, Zerecero (2021) estimates a model that includes a bias for living in one's birthplace and finds that it features smaller moving costs than a model that does not. This reflects an increase in the information the modeler has to predict future locations. The Shannon entropy, i.e. the average amount the modeler is surprised by any particular location choice, is smaller when they already know the person's birthplace. While it is a less direct comparison, Giannone et al. (2023) compares their estimated migration costs to the migration costs in Kennan and Walker (2011), and argues their new costs are lower because they include wealth in their model. This claim is consistent with wealth being an important piece of information about future location choices and the likelihood of moving.[24]

Other models also reduce the estimated moving cost by including features that help predict migration. For example, Heise and Porzio (2022) considers job search, where migration is more likely to occur conditional on a job offer, and Porcher (2020) considers rational inattention. Through the lens

---

[23]Given the differences across papers in currencies and years, it is not straightforward to compare the sizes, so I am guessing on a few of the comparisons.

[24]Giannone et al. (2023) and Kennan and Walker (2011) also differ in their geographies and timing, so the information effect of wealth is not the only difference between the two papers.

of my interpretation, prior to the decision to move, the modeler learns some information—either the agent gets a job offer (Heise and Porzio, 2022)[25] or they pick their optimal signal at a cost (Porcher, 2020). From the perspective of the modeler, this information helps predict where the agents are going to move to, lowering the Shannon entropy. Consistent with my interpretation, the estimated moving costs in these models are much lower.[26] Porcher (2020) estimates a model without his information frictions and finds the migration costs are 40 percent higher.

## 3.2   The role of random utility shocks

The key assumption of the standard model that generates my interpretation is that the random utility shocks are i.i.d. extreme value. In this section, I explain why that is the case and argue that relaxing that assumption could help estimate true moving costs.

If the standard moving cost model were the true model of the world, then migration costs could also be interpreted literally, in addition to the interpre-

---

[25]Schmutz and Sidibé (2019) has many similar features to Heise and Porzio (2022), but does not include a random utility shock. Through the lens of my framework, $\mu \to \infty$ in their model, meaning that the modeler has all the information they need to know whether a person will move or not, based on the job offer received and the other state variables. What this means is that moving costs only reflect the average increase in utility from moving, as in the formula in Appendix A.

[26]In Heise and Porzio (2022), I cannot directly apply the formulas in equations (2) and (3) because these models feature additional state variables for the agents. This means that the $v_{it} - v_{jt}$ term from equation (1) will not cancel out. The correct formulas for when there are state variables can be found in Appendix A, and these formulas include an additional term that represents the average utility gain from migration, net of moving costs and idiosyncratic shocks. I expect this additional term would be positive when migration also coincides with a job offer as in Heise and Porzio (2022). So this would actually lead to higher moving costs if there were no change in the modeler's information. Therefore, the decline in migration costs actually understates the improvement in the modeler's information.

tation I gave above. But in Section 1, I argued that arbitrary choices of the modeler played a large role in determining moving costs. If both of these things are true, it must be the standard moving cost model is somehow misspecified.[27]

Here, I argue that the source of this misspecification is the assumption that the preference shocks are i.i.d. This assumption is clearly unrealistic because it relies on there being no spatial correlation or persistence in the unobserved reasons that people choose locations.

To see this direct connection, suppose that the standard moving cost model was true, but that the researcher assumes that people draw i.i.d. shocks at the wrong frequency, i.e. every year rather than every five years. Then the researcher incorrectly assumes that in a five-year period, each agent has five draws in which they could get a high enough shock in order to move, whereas in truth, they only get one draw. When moving is rare, this means that for the same migration costs, the researcher would assume about five times more people move using the one-year model than the five-year model. So to match the data, they would infer higher moving costs for the five year model.

A similar argument can be made for subgeographies: if a researcher assumes agents get an independent shock for every MIGPUMA, then they will have to assume higher moving costs than if they assume one shock for the whole state in order to match the same amount of migration in the data.[28]

_____

[27]I do not mean this as a general criticism of moving cost models. After all, "all models are wrong, but some are useful" (Box, 1979). The point of this section is to focus on why this model is not that useful on one particular dimension—estimating moving costs—and to think about ways to change the assumptions such that it could.

[28]This point is acknowledged in Kennan and Walker (2011). However, they dismiss it by saying that the choice of whether to consider more preference shocks for populous states or lower moving costs does not affect their results on income, which was the original motivation of their paper.

However, if the researcher relaxed the i.i.d. assumption and instead assumed the shocks were highly correlated across time and space, then this argument would break down because the maximum of several random variables is not that different than the average, when the correlation is sufficiently high. A model with these correlations would estimate moving costs that are not as sensitive to the length of time periods or the size of geographies.

Allowing these correlations could also make the model less dependent on the information set of the modeler. For example, knowing the birthplace of an agent means that the modeler is able to predict that the agent will have a larger preference for living near their birthplace, in all time periods. Effectively, the birthplace is inducing a spatial and temporal correlation in the agents' match-specific preferences. If the correlation structure imposed by the modeler were rich enough, it might not matter if the modeler actually knew the agents' birthplaces or not.

Trying to get an estimate of "true" moving costs, that are not sensitive to timing and geographic assumptions, would require a model with both geographic and temporal correlation in location-person-match-specific preferences. This would be an interesting direction for future research.

# 4 Conclusion

Many people think of moving costs as a black box, since it is supposed to be a stand-in for many things that a modeler might not observe: information frictions, job and housing search, psychological costs of relocating, and, of

course, the actual monetary costs of moving. In this paper, I provide an alternative but related interpretation: average moving costs measure the size of the black box: moving costs are closely related to how little information is in the model about future locations.

Specifically, I show that the moving costs that are estimated by moving cost models measure the average Shannon entropy of next period's location minus the Shannon information of not moving. This measure is sensitive to seemingly-arbitrary choices of the economic modeler, and so I argue that it is not a good measure of actual moving costs. But what it does measure is nonetheless interesting, and helps us understand some of the various measures of moving costs in the literature.

# References

**Bartik, Alexander W and Kevin Rinz**, "Moving costs and worker adjustment to changes in labor demand: Evidence from longitudinal census data," 2018. Job Market Paper.

**Bayer, Christian and Falko Juessen**, "On the dynamics of interstate migration: Migration costs and self-selection," *Review of Economic Dynamics*, 2012, *15* (3), 377–401.

**Bishop, Kelly C**, "A dynamic model of location choice and hedonic valuation," 2012.

**Borusyak, Kirill, Rafael Dix-Carneiro, and Brian Kovak**, "Understanding migration responses to local shocks," 2022.

**Box, George E. P.**, "Robustness in the strategy of scientific model building," in Robert L. Launer and Geoffrey N. Wilkinson, eds., *Robustness in Statistics*, Academic Press, 1979, pp. 201–236.

**Bryan, Gharad and Melanie Morten**, "The aggregate productivity effects of internal migration: Evidence from Indonesia," *Journal of Political Economy*, 2019, *127* (5), 2229–2268.

_ , **Shyamal Chowdhury, and Ahmed Mushfiq Mobarak**, "Underinvestment in a profitable technology: The case of seasonal migration in Bangladesh," *Econometrica*, 2014, *82* (5), 1671–1748.

**Caliendo, Lorenzo, Maximiliano Dvorkin, and Fernando Parro**, "Trade and labor market dynamics: General equilibrium analysis of the china trade shock," *Econometrica*, 2019, *87* (3), 741–835.

**Fosgerau, Mogens and André de Palma**, "Generalized entropy models," 2016.

**Giannone, Elisa, Qi Li, Nuno Paixao, and Xinle Pang**, "Unpacking moving: A Spatial Equilibrium Model with Wealth," 2023.

**Heise, Sebastian and Tommaso Porzio**, "Labor Misallocation Across Firms and Borders," 2022.

**Howard, Greg and Hansen Shao**, "The Dynamics of Internal Migration: A New Fact and its Implications," 2022.

**Jia, Ning, Raven Molloy, Christopher Smith, and Abigail Wozniak**, "The economics of internal migration: Advances and policy questions," *Journal of Economic Literature*, 2023, *61* (1), 144–180.

**Jose, Victor Richmond R, Robert F Nau, and Robert L Winkler**, "Scoring rules, generalized entropy, and utility maximization," *Operations research*, 2008, *56* (5), 1146–1157.

**Kennan, John and James R Walker**, "The effect of expected income on individual migration decisions," *Econometrica*, 2011, *79* (1), 211–251.

**Koşar, Gizem, Tyler Ransom, and Wilbert van der Klaauw**, "Understanding Migration Aversion Using Elicited Counterfactual Choice Probabilities," *Journal of Econometrics*, 2022, *231* (1), 123–147. Annals Issue: Subjective Expectations & Probabilities in Economics.

**Matějka, Filip and Alisdair McKay**, "Rational inattention to discrete choices: A new foundation for the multinomial logit model," *American Economic Review*, 2015, *105* (1), 272–298.

**Oswald, Florian**, "The effect of homeownership on the option value of regional migration," *Quantitative Economics*, 2019, *10* (4), 1453–1493.

**Porcher, Charly**, "Migration with costly information," 2020. Job Market Paper.

**Ransom, Tyler**, "Labor market frictions and moving costs of the employed and unemployed," *Journal of Human Resources*, 2022, *57* (S), S137–S166.

**Ruggles, Steven, Sarah Flood, Matthew Sobek, Danika Brockman, Grace Cooper, Stephanie Richards, and Megan Schouweiler**, "IPUMS USA: Version 13.0 [dataset]," Online 2023.

**Schmutz, Benoît and Modibo Sidibé**, "Frictional labour mobility," *The Review of Economic Studies*, 2019, *86* (4), 1779–1826.

**Schubert, Gregor**, "House price contagion and U.S. city migration networks," 2021. Job Market Paper.

**Shannon, Claude E.**, "A Mathematical Theory of Communication," *The Bell System Technical Journal*, 1948, *27* (3), 379–423.

**Tombe, Trevor and Xiaodong Zhu**, "Trade, migration, and productivity: A quantitative analysis of china," *American Economic Review*, 2019, *109* (5), 1843–1872.

**Zerecero, Miguel**, "The Birthplace Premium," 2021. Job Market Paper.

# A   Extension to the model

In this appendix, I relax two big assumptions that I made in the main text: first, that agents are homogeneous except for their location; and second, that the model is in steady-state.

I derive a general formula for average moving costs, that is the sum of two terms. The first term is the same as in the main text of the paper. The second term is the average change in the continuation value $v$ for migrants, net of migration costs and idiosyncratic shocks.

I then consider a few special cases that are referenced in the main text. First, I consider the case where I drop the steady-state assumption but maintain the homogeneity assumption. I show that this new term is quantitatively small under reasonable assumptions on the symmetry of moving costs. Second, I consider the case where I maintain the steady-state assumption, but allow for the homogeneity assumption to be relaxed based on permanent characteristics of the agents. Finally, I consider the case with the steady-state assumption, but drop the homogeneity assumption to allow for consideration sets.

## A.1   More general setup

Consider an extension to the standard moving cost model, in which agents have a state $s$ that affects their payoffs and moving costs. $s$ is multidimensional, and it is a function of both the previous $s$, the location choice $i$, and a random variable $X$. This is a general setup so that $s$ could include age, the history of past locations, job status and wages, etc.

With the state variable, utility is now represented by this value function:

$$V_{nt}(j, s) = \max_i \log w_{it}(s) + a_{it}(s) - \delta_{ji}(s) + \frac{1}{\mu}\epsilon_{int} + \beta \mathbb{E} V_{nt+1}(i, s'(s, i, X))$$

where $w_{it}(s)$ is the (real) wage, $a_{it}(s)$ is the amenities in $i$, $\delta_{ji}(s)$ is the moving cost from $j$ to $i$, and $\epsilon_{int}$ is an i.i.d. extreme value shock. $\mu$ is a scale parameter, which governs the elasticity of substitution between places.

Define $v_{it}(j, s) \equiv \log w_{it}(s) + a_{it}(s) + \beta \mathbb{E} V_{nt+1}(i, s'(s, i, X))$. Then migration is given by

$$m_{j \rightarrow i,t}(s) = \frac{\exp(\mu(v_{it}(j, s) - \delta_{ji}(s)))}{\sum_k \exp(\mu(v_{kt}(j, s) - \delta_{jk}(s)))}$$

Again, I normalize $\delta_{ii}(s) = 0$, so the $\delta_{ji}(s)$ is then

$$\delta_{ji}(s) = v_{it}(j, s) - v_{jt}(j, s) - \frac{1}{\mu} \log m_{j \rightarrow i,t}(s) + \frac{1}{\mu} \log m_{j \rightarrow j,t}(s)$$

Consider the migration-weighted average moving cost in the steady-state of the model:

$$\begin{aligned}
\bar{\delta}_s &\equiv \frac{\sum_{s,i,j:i\neq j} p_i(s) m_{i \rightarrow j}(s) \delta_{ij}(s)}{\sum_{s,i,j:i\neq j} p_i(s) m_{i \rightarrow j}(s)} \\
&= \frac{1}{1 - \sum_{i,s} p_i(s) m_{i \rightarrow i}(s)} \sum_{s,i.j:i\neq j} p_i(s) m_{i \rightarrow j,t}(s) \left( -\frac{1}{\mu} \log m_{i \rightarrow j,t}(s) + \frac{1}{\mu} \log m_{i \rightarrow i,t}(s) \right) \\
&\quad + \frac{1}{\sum_{s,i,j:i\neq j} p_i(s) m_{i \rightarrow j}(s)} \sum_{s,i,j:i\neq j} p_i(s) m_{i \rightarrow j,t}(s) (v_{it}(j, s) - v_{jt}(j, s))
\end{aligned}$$

28

Rearranging,[29]

$$\bar{\delta}_s = \frac{1}{1 - \sum_i p_i(s) m_{i \to i}(s)} \frac{1}{\mu} \sum_i p_i(s) \left( - \sum_j [m_{i \to j}(s) \log m_{i \to j}(s)] + \log m_{i \to i}(s) \right)$$
$$+ \mathbb{E}_{ijs}^m (v_{it}(j, s) - v_{jt}(j, s))$$

Define $m_i \equiv \sum_{j:j \neq i} m_{i \to j}$ to be the total outmigration from $i$. Then

$$\bar{\delta}_s = \frac{1}{\mathbb{E}_{is} m_i(s)} \frac{1}{\mu} \mathbb{E}_{is} \left[ H(J|i, s) - I(i|i, s) \right] + \mathbb{E}_s^m [v_{it}(j, s) - v_{jt}(j, s)] \quad (5)$$

The first term is the same as before, except now the entropy and the informa-tion are both conditional on $s$. The $\delta$ is averaged across all $s$. The second term is the average gain in utility for migrants, net of moving costs and idiosyn-chratic shocks. The expectation $\mathbb{E}_s^m$ is the average, weighted by the number of migrants of type $s$ moving from $i$ to $j$.

In this more general setup, average moving costs are the sum of two compo-nents: the first is still a measure of the Shannon entropy from the perspective of the modeler minus the Shannon information of not-moving; and the second is the average gains from migration.

In the main text, the second term drops out because there are not average gains to migration. This is because the continuation value is the same for everyone, conditional on location, and I assumed the model was in steady-state, which meant that the same number of people moved into and out of each location.

---

[29]Note that $\sum_{j \neq i} m_{i \to j} = 1 - m_{i \to i}$.

## A.2 Special Case: Not in steady state

If I drop the steady-state assumption, but assume there are no additional $s$ states, then (5) becomes

$$\bar{\delta} = \frac{1}{\mathbb{E}_i m_i} \frac{1}{\mu} \mathbb{E}_i \left[ H(J|i) - I(i|i) \right] + \mathbb{E}^m [v_{it}(j) - v_{jt}(j)] \qquad (6)$$

The first term is the same as in the main part of the paper, and the second term is the additional utility gains from the fact that there is net migration to better places. The second term is likely to be positive, and it is small when differences in utility across space are small or when net migration is small.

In fact, I can numerically show that they are small, using the same data that I used in Section 2. I assume average moving costs into and out of every location are equal:

$$\sum m_{i \to j} \delta_{ij} = \sum m_{i \to j} \delta_{ji}$$

With this assumption, I can put a number on these average gains from migration net of moving costs and the idiosynchratic utility.[30] This assumption allows me to set up a system of two equations and two unknowns relating $\sum_{j \neq i} m_{i \to j}(v_{it} - v_{jt})$ and $\sum_{j \neq i} m_{i \to j} \delta_{ij}$, based on equation (1). Solving, the average gain from migration is given by:

$$\frac{1}{\mu} \sum_{i,j,i \neq j} p_i m_{i \to j} \log \left( \frac{m_{i \to j}}{m_{j \to i}} \frac{m_{j \to j}}{m_{i \to i}} \right)$$

---

[30]I cannot assume $\delta_{ij} = \delta_{ji}$ for every $i$ and $j$ because it overidentifies the data. With states, there would be $51 \times 51$ migration data points, but only 51 $v_i$'s and $\frac{51 \times 50}{2}$ $\delta_{ij}$'s to identify them with.

In the data, and with $\mu = 1$, this number is about 0.022. This is about 0.4 percent of the size of the information term (see Table 1). So at least in the standard model, the steady-state assumption was not quantitatively affecting my results.

## A.3  Special Case: $s$ is immutable

Another special case of the more general result is if $s$ is immutable: i.e. $s' = s$, and I maintain the steady-state assumption. $s$ being immutable means I can rewrite $v_{it}(j, s) \equiv v_{its}$, i.e. the continuation value does not depend on $j$. And because the model is in steady-state, the number of people of type $s$ moving into $i$ is cancelled out by the number of people of type $s$ moving out of $i$. So the average gains from migration term drops out:

$$\bar{\delta} = \frac{1}{\mathbb{E}_{is} m_i(s)} \frac{1}{\mu} \mathbb{E}_{is} \left[ H(J|i, s) - I(i|i, s) \right]$$

This leaves us with the main result again, but where the Shannon entropy and the Shannon information are conditional on $s$.

## A.4  $s$ and $j$ do not affect $v_{it}$

Another straightforward example is if $v_{it}(j, s)$ does not depend on $j$ or $s$, e.g. $s$ governs the contemporaneous moving costs, but nothing else.[31] This could be the case if, at the start of each period, each agent drew a random consideration set, which is represented by $s$. But once they moved to the new region, they

---

[31]As in the last example, I maintain the steady-state assumption here.

looked just like anyone else there.

Under this assumption, we again get the same equation:

$$\bar{\delta} = \frac{1}{\mathbb{E}_{is} m_i(s)} \frac{1}{\mu} \mathbb{E}_{is} \left[ H(J|i,s) - I(i|i,s) \right]$$

So the Shannon entropy and Shannon information depend on $s$, but the result is otherwise the same.

In general, however, adding the state to the model leads to the possibility that there are average gains to in utility for migrants. Examples of $s$ that would matter are if $s$ equals the history of past locations, age, or employment status. I would expect these to be positive because migrants will tend to move to places with higher utility for themselves.